

WeightWatcher: A Diagnostic Tool for Deep Neural Networks

Charles H. Martin PhD, Calculation Consulting

(in joint with Michael Mahoney, UC Berkeley)
`pip install weightwatcher`

Open source tool: weightwatcher

<https://github.com/CalculatedContent/WeightWatcher>

WeightWatcher (WW): is an open-source, diagnostic tool for analyzing Deep Neural Networks (DNN), without needing access to training or even test data. It can be used to:


- analyze pre/trained pyTorch and keras models
- inspect models that are difficult to train
- gauge improvements in model performance
- predict test accuracies across different models
- detect potential problems when compressing or fine-tuning pretrained models

It is based on theoretical research (done in joint with UC Berkeley) into *Why Deep Learning Works*, using ideas from Random Matrix Theory (RMT), Statistical Mechanics, and Strongly Correlated Systems.

pip install weightwatcher

Shape and Scale Metrics

WeightWatcher (WW): analyzes the shape and scale of the correlations in the layer weight matrices:



WW: extracts, plots, and fits the Empirical Spectral Density (ESD, or eigenvalues) for each layer weight matrix (or tensor slice).

The *tail of the ESD* contains the most informative components.

The shape of the tail carries useful information!

WeightWatcher: Usage

Usage

```
import weightwatcher as ww
watcher = ww.WeightWatcher(model=model)
details = watcher.analyze(plot=True)
summary = watcher.get_summary(details)
```

	layer_id	name	D	M	N	alpha	alpha_weighted	has_esd	lambda_max	layer_type	rand_num_spikes	rand_sigma_mp	ri
0	2	None	0.240111	3.0	64.0	2.400712	2.627967	1.0	12.435451	LAYER_TYPE.CONV2D	...	136.0	1.000000
1	5	None	0.112669	64.0	128.0	7.116304	4.721276	1.0	4.607285	LAYER_TYPE.CONV2D	...	25.0	0.551250
2	8	None	0.076209	128.0	256.0	2.981087	1.739893	1.0	3.833927	LAYER_TYPE.CONV2D	...	17.0	0.451562
3	10	None	0.068890	256.0	256.0	5.667264	2.600458	1.0	2.876445	LAYER_TYPE.CONV2D	...	0.0	0.935068
4	13	None	0.084938	256.0	512.0	2.593428	1.432684	1.0	3.568032	LAYER_TYPE.CONV2D	...	8.0	0.431523
5	15	None	0.038416	512.0	512.0	3.309962	2.216486	1.0	4.673487	LAYER_TYPE.CONV2D	...	0.0	0.939111
6	18	None	0.052924	512.0	512.0	3.446656	1.859810	1.0	3.464163	LAYER_TYPE.CONV2D	...	0.0	0.886574
7	20	None	0.034290	512.0	512.0	3.261262	2.524426	1.0	5.943799	LAYER_TYPE.CONV2D	...	0.0	0.942012
8	25	None	0.032563	4096.0	25088.0	2.325065	3.583809	1.0	34.784030	LAYER_TYPE.DENSE	...	1.0	0.896506
9	28	None	0.030081	4096.0	4096.0	2.167513	3.856526	1.0	60.278519	LAYER_TYPE.DENSE	...	1.0	0.959863
10	31	None	0.039773	1000.0	4096.0	2.825653	4.999373	1.0	58.786867	LAYER_TYPE.DENSE	...	206.0	1.000000

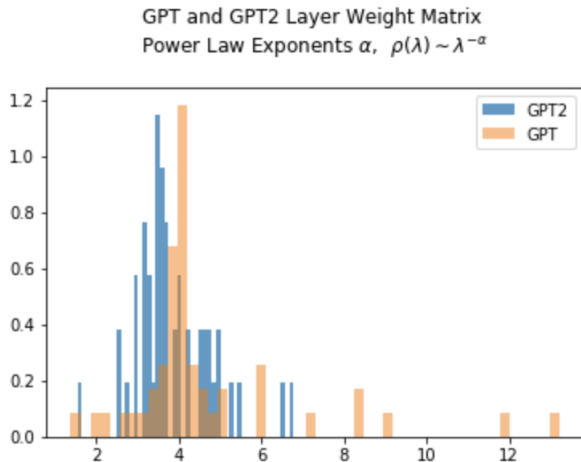
```
summary = watcher.get_summary(details)
```

```
{'log_norm': 2.11,  
 'alpha': 3.06,  
 'alpha_weighted': 2.78,  
 'log_alpha_norm': 3.21,  
 'log_spectral_norm': 0.89,  
 'stable_rank': 20.90,  
 'mp_softrank': 0.52}]
```

Layer-by-Layer Analysis

WW layer metrics: can detect potential problems *in the ESD shapes*

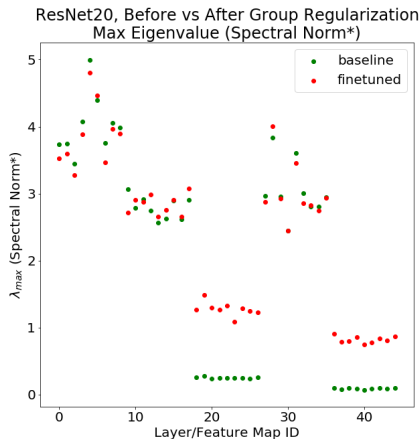
Poorly trained models (**orange**) can have unusually large layer α 's.



Layer-by-Layer Analysis

WW layer metrics: can detect potential problems *in the ESD Scales*


Compressed models (red) can show unexpected scale changes



example from Intel distiller Group Regularization technique


α : a regularization metric

The WW $\langle\alpha\rangle$ metric: predicts test accuracy for a given model (i.e same depth) when varying the regularization hyper-parameters (such as batch size, weight decay, momentum, etc.)—*without access to the test or training data*.



$\hat{\alpha}$: a multi-purpose metric

The WW $\hat{\alpha}$ metric: predicts test accuracy for models in the same architecture series across varying depth and other architecture parameters and regularization hyper-parameters—*without access to the test or training data*.



Research Update: Early Stopping with α

Early Stopping: Our HTSR theory suggests that when our PL metric $\alpha < 2.0$, the layer may be overtrained. Moreover, the early stopping should be applied when the layer averaged $\langle \alpha \rangle < 2.0$.



Early results by Xander Dunn show this amazingly well!

Here, when training a transformer model, as the training error decreases, α decreases. But just when the test error starts to increase, $\alpha < 2.0$.

This is remarkable, and we are pursuing this on a wide variety of models.

Research Update: Early Stopping with α

Early Stopping: Our HTSR theory suggests that when our PL metric $\alpha < 2.0$, the layer may be overtrained. Moreover, the early stopping should be applied when the layer averaged $\langle \alpha \rangle < 2.0$.




Early results by Xander Dunn show this amazingly well!

Here, when training a transformer model, as the training error decreases, α decreases. But just when the test error starts to increase, $\alpha < 2.0$.


This is remarkable, and we are pursuing this on a wide variety of models.

Research Update: Heavy Tails in \mathbf{W} and \mathbf{X}

Correlation Traps We *conjecture* that when the unusually large elements $W_{i,j}$ arise in the weight matrices, these act like traps (in the ESD of \mathbf{X} that prevent good generalization.



(a) ESD of \mathbf{W} and randomized \mathbf{W} .



(b) ESD of \mathbf{W} and randomized \mathbf{W} .

These can be seen using the [randomize] option in **weightwatcher**.
We believe we can use this to detect overfitting and are investigating this.